

High-Dimensional Low-Rank Tensor Autoregressive Time Series Modeling

Presenter: Zhi-Long Han

December 29, 2025

Reference: Di Wang, Yao Zheng, Guodong Li, *High-Dimensional Low-Rank Tensor Autoregressive Time Series Modeling*, arXiv:2101.04276, 2021.

Outline

1. Background & Motivation
2. Model & Problem Setting
3. Estimation Methods
4. Algorithms
5. Theory Highlights
6. Experimental Design
7. Simulation Studies
8. Real Data Analysis
9. Discussion & Takeaways
10. References

Why Tensor-Valued Time Series?

Tensor-valued observation: at each time t we observe

$$Y_t \in \mathbb{R}^{p_1 \times p_2 \times \cdots \times p_d}, \quad E_t \in \mathbb{R}^{p_1 \times \cdots \times p_d}.$$

- ▶ Examples: portfolio panels, multi-way traffic flows, dynamic images/videos.
- ▶ Tensor keeps **multi-way structure** that vectorization destroys.

Naïve baseline: vectorize then VAR(1)

$$y_t = Ay_{t-1} + e_t, \quad y_t = \text{vec}(Y_t) \in \mathbb{R}^p, \quad p = \prod_{i=1}^d p_i, \quad A \in \mathbb{R}^{p \times p}.$$

- ▶ **Parameter explosion:** p^2 entries in $A \Rightarrow$ infeasible when p is large and T is limited.
- ▶ **Interpretability loss:** no longer see effects along each mode.

Motivation: What Structure Can We Exploit?

Key Insight

Replace the huge VAR matrix $A \in \mathbb{R}^{p \times p}$ by a **transition tensor**

$$\mathcal{A} \in \mathbb{R}^{p_1 \times \cdots \times p_d \times p_1 \times \cdots \times p_d}$$

with **low Tucker multilinear ranks** $r = (r_1, \dots, r_{2d})$, where $r_k \ll p_k$.

What this paper delivers

- ▶ LRTAR model + dynamic factor interpretation via Tucker.
- ▶ Two estimation regimes: **LTR** (asymptotic) and **SSN/TSSN** (high-dimensional).
- ▶ New SSN regularizer (all square matricizations) + rank selection consistency.

LRTAR(1): Model and VAR View

Tensor autoregression:

$$Y_t = \langle \mathcal{A}, Y_{t-1} \rangle + E_t, \quad (1)$$

where

- ▶ $Y_t, E_t \in \mathbb{R}^{p_1 \times \cdots \times p_d}$,
- ▶ $\mathcal{A} \in \mathbb{R}^{p_1 \times \cdots \times p_d \times p_1 \times \cdots \times p_d}$ is a $2d$ -th order transition tensor,
- ▶ contraction over the last d modes:

$$(\langle \mathcal{A}, Y_{t-1} \rangle)_{i_1, \dots, i_d} = \sum_{j_1=1}^{p_1} \cdots \sum_{j_d=1}^{p_d} \mathcal{A}_{i_1, \dots, i_d, j_1, \dots, j_d} (Y_{t-1})_{j_1, \dots, j_d}.$$

VAR view (square matricization): let $p = \prod_{i=1}^d p_i$, $y_t = \text{vec}(Y_t) \in \mathbb{R}^p$. Define $\mathcal{A}_{[S_2]} \in \mathbb{R}^{p \times p}$ by grouping (i_1, \dots, i_d) as rows and (j_1, \dots, j_d) as columns. Then

$$y_t = \mathcal{A}_{[S_2]} y_{t-1} + e_t, \quad e_t = \text{vec}(E_t),$$

where $\rho(\mathcal{A}_{[S_2]}) < 1$.

Low-Rank Transition Tensor

Low-rank structure assumption (Tucker):

$$\mathcal{A} = \mathcal{G} \times_1 U_1 \times_2 U_2 \cdots \times_{2d} U_{2d}, \quad (2)$$

where

- ▶ $\mathcal{A} \in \mathbb{R}^{p_1 \times \cdots \times p_d \times p_1 \times \cdots \times p_d}$
- ▶ Core tensor: $\mathcal{G} \in \mathbb{R}^{r_1 \times \cdots \times r_{2d}}$
- ▶ Factor matrices: $U_k \in \mathbb{R}^{p_k \times r_k}$ ($k = 1, \dots, d$).

Why Tucker helps (parameter reduction):

Model	Number of parameters (order)
Naïve VAR	p^2 ($p = \prod_{k=1}^d p_k$)
LRTAR with Tucker	$\prod_{k=1}^{2d} r_k + \sum_{k=1}^d p_k r_k + \sum_{k=1}^d p_k r_{d+k}$

Dynamic Factor Interpretation

low-dimensional dynamics


Let $\mathcal{F}_{t-1} = Y_{t-1} \times_1 U_1^\top \times_2 \cdots \times_d U_d^\top \in \mathbb{R}^{r_1 \times \cdots \times r_d}$. Then the response-side projection satisfies

$$Y_t \times_1 U_{d+1}^\top \times_2 \cdots \times_d U_{2d}^\top = \langle \mathcal{G}, \mathcal{F}_{t-1} \rangle + E_t \times_1 U_{d+1}^\top \times_2 \cdots \times_d U_{2d}^\top. \quad (3)$$

Hence, (3) is a **low-dim tensor regression** defined on the projections of Y_t and Y_{t-1} .

Interpretation

- ▶ U_1, \dots, U_d : **predictor loadings** (compress Y_{t-1}).
- ▶ U_{d+1}, \dots, U_{2d} : **response loadings** (compress Y_t).
- ▶ \mathcal{G} links low-dim predictors to low-dim responses.



Two Estimation Regimes

Regime A: Low-dimensional (exact low-rank)

- ▶ Tucker ranks r are known; T is large relative to effective dimension.
- ▶ Estimate by **rank-constrained least squares**.
- ▶ Result: **asymptotic normality** (Theorem 1).

Regime B: High-dimensional (approximate low-rank)

- ▶ $p = \prod p_i$ large; \mathcal{A} is only *approximately* low-rank.
- ▶ Estimate by convex regularization:
 - ▶ SN: sum of nuclear norms of one-mode matricizations.
 - ▶ MN: nuclear norm of a single $p \times p$ square matricization.
 - ▶ **SSN (proposed):** sum of nuclear norms of *all* square matricizations.
- ▶ Result: **non-asymptotic error bounds + rank selection** (Theorems 2–5).

Low-Dimensional Estimator: low-Tucker-rank (LTR)

Objective (LTR):

$$\begin{aligned}\widehat{\mathcal{A}}_{\text{LTR}} = \arg \min_{\mathcal{A}} \frac{1}{T-1} \sum_{t=2}^T & \left\| Y_t - \langle \mathcal{A}, Y_{t-1} \rangle \right\|_F^2 \\ \text{s.t. } \text{rank}_k(\mathcal{A}) \leq r_k, \quad k = 1, \dots, 2d.\end{aligned}\tag{4}$$

Unknowns and sizes:

- ▶ Transition tensor: $\mathcal{A} \in \mathbb{R}^{p_1 \times \dots \times p_d \times p_1 \times \dots \times p_d}$.
- ▶ Tucker factors and core:

$$\begin{aligned}U_k \in \mathbb{R}^{p_k \times r_k} \quad (k = 1, \dots, d), \quad U_{d+k} \in \mathbb{R}^{p_k \times r_{d+k}} \quad (k = 1, \dots, d), \\ \mathcal{G} \in \mathbb{R}^{r_1 \times \dots \times r_{2d}}, \quad \mathcal{A} = \mathcal{G} \times_1 U_1 \cdots \times_{2d} U_{2d}.\end{aligned}$$

Computation

Solve (4) via **ALS** (alternating LS / Procrustes updates for U_k and \mathcal{G}).

Algorithm 1: ALS for LTR

Algorithm: Alternating Least Squares (ALS)

Inputs $\{Y_t\}_{t=1}^T$, ranks $r = (r_1, \dots, r_{2d})$, max iters S , tolerance ε

Outputs $\widehat{\mathcal{A}}_{\text{LTR}}$ (or $\widehat{\mathcal{G}}, \{\widehat{U}_k\}_{k=1}^{2d}$)

Init Compute $\mathcal{A}^{(0)}$ (e.g., OLS/RRR), then HOSVD to obtain $\mathcal{G}^{(0)}, \{U_k^{(0)}\}_{k=1}^{2d}$

Updating (1) Update predictor loadings $U_k^{(s+1)} \in \mathbb{R}^{p_k \times r_k}, k = 1, \dots, d$ (*Procrustes*)
(2) Update response loadings $U_{d+k}^{(s+1)} \in \mathbb{R}^{p_k \times r_{d+k}}, k = 1, \dots, d$ (*Procrustes*)
(3) Update core $\mathcal{G}^{(s+1)} \in \mathbb{R}^{r_1 \times \dots \times r_{2d}}$ (*LS*)
(4) Form $\mathcal{A}^{(s+1)} = \mathcal{G}^{(s+1)} \times_1 U_1^{(s+1)} \dots \times_{2d} U_{2d}^{(s+1)}$

Stop If $\|\mathcal{A}^{(s+1)} - \mathcal{A}^{(s)}\|_F / \|\mathcal{A}^{(s)}\|_F \leq \varepsilon$

Theorem 1: Asymptotic Normality of LTR (Low-dimensional)

Theorem 1

Let $\hat{\mathcal{A}}_{LTR}$ be the rank-constrained LS estimator. Assume $\mathbb{E}\|e_t\|_2^2 < \infty$ and stationarity holds. Then

$$\sqrt{T} \operatorname{vec}((\hat{\mathcal{A}}_{LTR} - \mathcal{A})_{[S_2]}) \Rightarrow \mathcal{N}(0, \Sigma_{LTR}),$$

where $(\cdot)_{[S_2]}$ is the square matricization and $\Sigma_{LTR} = H(H^\top J H)^\dagger H^\top \in \mathbb{R}^{p^2 \times p^2}$.

Interpretation:

If the Tucker ranks are correctly specified, LTR is a **well-specified** low-dimensional estimator; under Gaussian noise, least squares coincides with the **Maximum Likelihood Estimator (MLE)**, yielding the classical $\mathcal{O}(T^{-1/2})$ estimation rate.

High-Dimensional Regularization: Unified Template

When p is large, estimate \mathcal{A} via:

$$\hat{\mathcal{A}} = \arg \min_{\mathcal{A}} \frac{1}{T-1} \sum_{t=2}^T \|Y_t - \langle \mathcal{A}, Y_{t-1} \rangle\|_F^2 + \lambda \mathcal{R}(\mathcal{A}). \quad (5)$$

Approximate low-rank matrix class For a matrix $M \in \mathbb{R}^{m \times n}$ with singular values $\sigma_j(M)$, define

$$\mathbb{B}_q(r_q; m, n) = \left\{ M : \sum_{j \geq 1} \sigma_j(M)^q \leq r_q, \ 0 \leq q < 1 \right\}.$$

- ▶ $q = 0$: exactly rank- r_0 matrices.
- ▶ $0 < q < 1$: singular values decay (nearly low-rank).

SN vs MN vs SSN: Regularizers

Abbrev.: SN: Sum of Nuclear norms; MN: Matrix Nuclear norm;
SSN: Sum of Square Nuclear norms.

Notation: $p_{-i} = \prod_{j=1, j \neq i}^d p_j$, $p = \prod_{j=1}^d p_j$.

Method	Matrix shape (how obtained + size)	Regularizer
SN	Matricize along mode i , $i = 1, \dots, 2d$ $\mathcal{A}_{(i)} \in \mathbb{R}^{p_i \times (p_{-i}p)}$	$\mathcal{R}_{SN}(\mathcal{A}) = \sum_{i=1}^{2d} \ \mathcal{A}_{(i)}\ _*$
MN	VAR view; one fixed square view S_1 $\mathcal{A}_{[S_1]} \in \mathbb{R}^{p \times p}$	$\mathcal{R}_{MN}(\mathcal{A}) = \ \mathcal{A}_{[S_1]}\ _*$
SSN	All square views I_k , $k = 1, \dots, 2^{d-1}$ $\mathcal{A}_{[I_k]} \in \mathbb{R}^{p \times p}$	$\mathcal{R}_{SSN}(\mathcal{A}) = \sum_{k=1}^{2^{d-1}} \ \mathcal{A}_{[I_k]}\ _*$

Takeaway

SSN keeps **square-matrix structure** (like MN) but enforces it across *all* square views.

Algorithm 2: ADMM for SSN

To handle multiple nuclear norms in $\|\mathcal{A}\|_{SSN}$, introduce auxiliary variables $W_k \in \mathbb{R}^{p_1 \times \dots \times p_d \times p_1 \times \dots \times p_d}$ ($k = 1, \dots, 2^{d-1}$) with constraints $W_k = \mathcal{A}$.

Augmented Lagrangian Function

$$\mathcal{L}(\mathcal{A}, \{W_k\}, \{C_k\}) = L_T(\mathcal{A}) + \sum_{k=1}^{2^{d-1}} \left(\lambda_{SSN} \|(W_k)_{[I_k]}\|_* + \frac{\rho}{2} \|\mathcal{A} - W_k + C_k\|_F^2 \right).$$

ADMM iterations (Algorithm 2)

- ▶ **A-update:** $\mathcal{A}^{(j+1)} = \arg \min_{\mathcal{A}} L_T(\mathcal{A}) + \frac{\rho}{2} \sum_k \|\mathcal{A} - W_k^{(j)} + C_k^{(j)}\|_F^2$.
- ▶ **W-update (parallel over k):**

$$(W_k^{(j+1)})_{[I_k]} = \text{SVT}_{\lambda_{SSN}/\rho} \left((\mathcal{A}^{(j+1)} + C_k^{(j)})_{[I_k]} \right),$$

where $\text{SVT}_\tau(\cdot)$ soft-thresholds singular values by τ .

- ▶ **Dual:** $C_k^{(j+1)} = C_k^{(j)} + \mathcal{A}^{(j+1)} - W_k^{(j+1)}$.

Assumptions and Key Constants

High-level assumptions

- ▶ Stationary VAR(1) representation (spectral radius < 1).
- ▶ Innovations are sub-Gaussian with parameter κ (controls tails).
- ▶ Restricted strong convexity (RSC): loss $L_T(\mathcal{A})$ behaves like a quadratic over low-rank tangent cones.

Notation used in bounds Let $\Sigma_e = \text{Var}(e_t) \in \mathbb{R}^{p \times p}$, and define constants

$$M_1 = \frac{\lambda_{\max}(\Sigma_e)}{\mu_{\min}(\mathcal{A})^{1/2}}, \quad M_2 = \frac{\lambda_{\min}(\Sigma_e) \mu_{\max}(\mathcal{A})}{\lambda_{\max}(\Sigma_e) \mu_{\min}(\mathcal{A})},$$

where μ_{\min}, μ_{\max} capture stability/conditioning of the process.

Role

Bad stability / ill-conditioning \Rightarrow small $M_2 \Rightarrow$ larger sample size needed.

Theorem 2: SN Estimator (Non-asymptotic Error Bound)

SN estimator

$$\hat{\mathcal{A}}_{SN} = \arg \min_{\mathcal{A}} L_T(\mathcal{A}) + \lambda_{SN} \sum_{i=1}^{2d} \|\mathcal{A}_{(i)}\|_*.$$

Theorem 2 (core condition + bound)

If $T \gtrsim \max_{1 \leq i \leq d}(p_{-i}p) + \max(\kappa^2, \kappa^4)M_2^{-2}p$, $\lambda_{SN} \gtrsim \kappa^2 M_1 d^{-2} \sum_{i=1}^{2d} \sqrt{\frac{p_{-i}p}{T}}$, and $\mathcal{A}_{(i)} \in \mathbb{B}_q(r_q^{(i)}; p_i, p_{-i}p)$ for $0 \leq q < 1$, then w.h.p.

$$\|\hat{\mathcal{A}}_{SN} - \mathcal{A}\|_F \lesssim \sqrt{r_q} \left(\frac{2d \lambda_{SN}}{\alpha_{RSC}} \right)^{1-q/2}, \quad r_q = \frac{1}{2d} \sum_{i=1}^{2d} r_q^{(i)}.$$

Why worse: $\mathcal{A}_{(i)} \in \mathbb{R}^{p_i \times (p_{-i}p)}$ is **highly unbalanced** \Rightarrow deviation dominated by $p_{-i}p$.

Theorem 3–4: MN vs SSN

Theorem 3 (MN) and Theorem 4 (SSN): bounds

MN: if $T \gtrsim [1 + \max(\kappa^2, \kappa^4)M_2^{-2}]p$ and $\lambda_{MN} \gtrsim \kappa^2 M_1 \sqrt{p/T}$, with
 $\mathcal{A}_{[S_1]} \in \mathbb{B}_q(s_q^{(1)}; p, p)$,

$$\|\hat{\mathcal{A}}_{MN} - \mathcal{A}\|_F \lesssim \sqrt{s_q^{(1)}} \left(\frac{\lambda_{MN}}{\alpha_{RSC}} \right)^{1-q/2}.$$

SSN: if $T \gtrsim [1 + \max(\kappa^2, \kappa^4)M_2^{-2}]p$ and $\lambda_{SSN} \gtrsim \kappa^2 M_2^{1-d} \sqrt{p/T}$, with
 $\mathcal{A}_{[I_k]} \in \mathbb{B}_q(s_q^{(k)}; p, p)$,

$$\|\hat{\mathcal{A}}_{SSN} - \mathcal{A}\|_F \lesssim \sqrt{s_q} \left(\frac{2^{d-1} \lambda_{SSN}}{\alpha_{RSC}} \right)^{1-q/2}, \quad s_q = 2^{1-d} \sum_{k=1}^{2^{d-1}} s_q^{(k)}.$$

Sample Size & Error Rate Comparison

Define $p_{-i} = \prod_{j \neq i} p_j$, $p = \prod_{j=1}^d p_j$.

Estimator	Sample size condition (order)	Frobenius error rate (order)
SN	$T \gtrsim \max_{i \leq d}(p_{-i}p) + M_2^{-2}p$	$\sqrt{r_q} \left(\max_{i \leq d} \frac{p_{-i}p}{T} \right)^{\frac{1}{2} - \frac{q}{4}}$
MN	$T \gtrsim (1 + M_2^{-2})p$	$\sqrt{s_q^{(1)}} \left(\frac{p}{T} \right)^{\frac{1}{2} - \frac{q}{4}}$
SSN	$T \gtrsim (1 + M_2^{-2})p$	$\sqrt{s_q} \left(\frac{p}{T} \right)^{\frac{1}{2} - \frac{q}{4}}$

How to read this table

- ▶ **SN** is slower because $p_{-i}p$ can be much larger than p (unbalanced unfoldings).
- ▶ MN improves by using a **single** $p \times p$ square view.
- ▶ **SSN** keeps the same p/T scaling but is **more robust** by enforcing low-rank across *all* square views.

Theorem 5: TSSN Rank Selection Consistency

Goal: consistently recover the true Tucker ranks $\{\text{rank}(\mathcal{A}_{(i)})\}_{i=1}^{2d}$. **Mechanism:**

Truncated SSN (TSSN)

- ▶ Obtain $\widehat{\mathcal{A}}_{SSN}$ first, then perform **Hard Thresholding** on singular values of each unfolding $(\widehat{\mathcal{A}}_{SSN})_{(i)}$.
- ▶ Keep singular values $\sigma_j \geq \gamma$ (typically $\gamma \propto \lambda_{SSN}$), discard the rest.

Theorem 5 (statement)

Under the conditions of Theorem 4 and an additional signal strength assumption (Assumption 3), if $\lambda_{SSN} \asymp \kappa^2 M_2^{1-d} \sqrt{p/T}$, then as $T \rightarrow \infty$ (fixed d),

$$\mathbb{P}\left(\text{rank}((\widehat{\mathcal{A}}_{TSSN})_{(i)}) = \text{rank}(\mathcal{A}_{(i)}), i = 1, \dots, 2d\right) \rightarrow 1,$$

and

$$\|\widehat{\mathcal{A}}_{TSSN} - \mathcal{A}\|_F = \mathcal{O}_p\left(\sqrt{s_0 \frac{p}{T}}\right),$$

where s_0 is the exact low-rank complexity parameter in Assumption 3.

Experimental Design: Two-Track Validation

Track 1: Simulation Studies (Validation)

- ▶ **Goal:** Verify estimation accuracy and validate theoretical scaling rates.
- ▶ **Scope:** Low-Dim (LTR vs RRR) & High-Dim (SN vs MN vs SSN).
- ▶ **Metric:** Tensor estimation error $\|\hat{\mathcal{A}} - \mathcal{A}\|_F$.

Track 2: Real Data Analysis (Application)

- ▶ **Goal:** Test predictive power and interpret learned low-rank structures.
- ▶ **Datasets:** 10×10 Matrix Portfolios ($d = 2$) & $4 \times 4 \times 2$ Tensor Portfolios ($d = 3$).
- ▶ **Metric:** Rolling forecasting error (RMSFE).

Benchmarking: Baselines & Estimators

1. Existing Baseline Methods

- ▶ Vector-based: VAR (Standard Vector AR), VFM (Vector Factor Model).
- ▶ Tensor-based: MAR/MTAR (Separable), MFM/TFM (Factor models).

2. Proposed LRTAR Estimators (Ours)

Low-Dimensional Regime	High-Dimensional Regime
LTR (Rank-constrained LS)	SN / MN (Standard norms) SSN (Sum of Square-matrix Nuclear norms) TSSN (SSN + Truncation)

Simulation Setup & Protocols

1. Data Generation (Controlled Environment)

- ▶ **Process:** Generate stationary tensor series from LRTAR(1):

$$\mathfrak{Y}_t = \langle \mathcal{A}, \mathfrak{Y}_{t-1} \rangle + \mathcal{E}_t, \quad \mathcal{E}_t \sim \mathcal{N}(0, I).$$

- ▶ **Variables:** Sample size T , Dimension $p = \prod p_i$, Ranks r , Noise levels.

2. Estimators Comparison Groups

- ▶ **Low-Dimensional Regime:**

[LTR](#) (Ours) vs. [RRR](#) vs. [OLS](#) — *Test efficiency gain from tensor structure.*

- ▶ **High-Dimensional Regime:**


[SN](#) vs. [MN](#) vs. [SSN](#) vs. [TSSN](#) — *Test benefits of square matricization.*

Simulation I: Low-Dimensional Estimation Efficiency

Setup: Low-rank LRTAR model ($d = 2, 3$) with varying dimensions p_i and ranks r .

Key Finding

The proposed **LTR estimator** achieves the lowest estimation error by strictly enforcing tensor rank constraints, significantly outperforming RRR and OLS.




Simulation II: Verification of Theoretical Rates (High-Dim)

Setup: Verify error bound for SSN estimator. Varying p , T , and complexity s_0 .

Key Finding

The estimation error scales **linearly** with total dimension p and complexity s_0 , confirming that SSN is robust and effective for high-dimensional scaling.




Simulation III: Comparison of High-Dimensional Estimators

Comparison: SN vs. MN vs. SSN vs. TSSN.

Key Finding

TSSN yields the lowest error by combining the efficiency of balanced square matricizations (SSN) with rank truncation, significantly beating SN and MN.



Real Data: Fama-French Portfolios Analysis


Data: Monthly returns of 10×10 portfolios ($d = 2$) & $4 \times 4 \times 2$ portfolios ($d = 3$).

Metric: Out-of-sample rolling forecast error.

Key Results

- ▶ **Forecasting:** LRTAR (TSSN) achieves the lowest out-of-sample error.
- ▶ **Structure:** Identified factors show *Market* (uniform) and *Size/Value* (monotonic) patterns.

Model	VAR	VFM	MAR	MFM	LRTAR		Best	Worst	
					SSN	TSSN			
In-sample	ℓ_2 norm	31.61	35.85	34.12	35.86	33.18	33.38	VAR	MFM
	ℓ_0 norm	8.09	9.23	9.12	9.24	8.84	8.89	VAR	MFM
Out-of-sample	ℓ_2 norm	39.84	39.11	35.26	38.53	32.60	31.47	TSSN	VAR
	ℓ_∞ norm	11.98	12.30	11.47	11.00	9.53	9.33	TSSN	VFM



Experimental Summary & Implications

Simulation Insights

- ▶ **Low-Dim:** LTR beats vector baselines when ranks are specified correctly.
- ▶ **High-Dim:**
 - ▶ SN suffers from unbalanced (rectangular) matricizations.
 - ▶ MN/SSN gain efficiency via balanced square views.
 - ▶ TSSN further improves accuracy via rank truncation.

Real Data Evidence

- ▶ **Forecasting:** SSN/TSSN achieve the best out-of-sample performance.
- ▶ **Interpretability:**
 - ▶ Response-side low ranks suggest a *low-dimensional predictable subspace*.
 - ▶ Factors align with financial theory (Market/Size effects).

Conclusion

Structured low-rank dynamics (LRTAR) significantly improve both **prediction accuracy** and **model interpretability**.

Main Takeaways

1. The Methodology: LRTAR

$$\mathfrak{Y}_t = \langle \mathcal{A}, \mathfrak{Y}_{t-1} \rangle + \mathcal{E}_t, \quad \mathcal{A} \text{ has low Tucker-rank.}$$

2. Key Contributions

- ▶ **Two Regimes:** Asymptotic normality (Low-dim) vs. Finite-sample rates (High-dim).
- ▶ **SSN Regularizer:** The key innovation. Exploits **balanced square matricizations** to beat standard SN.
- ▶ **TSSN Estimator:** Provides rank recovery with theoretical consistency.

Limitations & Future Directions

- ▶ Currently limited to AR(1); extending to AR(p) requires stability selection.
- ▶ Computational cost: SSN-ADMM involves 2^{d-1} SVDs (heavy computation when total dimension p is very large).

References

- ▶ ([Wang et al.'24](#)) "High-dimensional low-rank tensor autoregressive time series modeling." *Journal of Econometrics*, 2024, 238(1): 105544.
- ▶ ([Wang et al.'22](#)) "High-dimensional vector autoregressive time series modeling via tensor decomposition." *Journal of the American Statistical Association*, 2022, 117(539): 1338–1356.
- ▶ ([Mu et al.'14](#)) "Square Deal: Lower Bounds and Improved Relaxations for Tensor Recovery." *ICML*, 2014.
- ▶ ([Negahban & Wainwright'11](#)) "Estimation of (near) low-rank matrices with noise and high-dimensional scaling." *Annals of Statistics*, 2011.
- ▶ ([Lam et al.'12](#)) "Factor modeling for high-dimensional time series: inference for the number of factors." *Annals of Statistics*, 2012.